

# AmazonFineFoodReviews\_DT

September 2, 2018

## 1 Decision Tree Amazon Fine Food Reviews

In [15]: %matplotlib inline

```
import sqlite3
import pandas as pd #for data frames
import numpy as np #numpy array operations
import nltk #natural lang processing, for processing text
import string
import matplotlib.pyplot as plt
import seaborn as sns #for plotting
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer

from nltk.stem.porter import PorterStemmer
import pickle
import seaborn as sn

import matplotlib.pyplot as plt
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
from sklearn.model_selection import train_test_split
from sklearn.metrics import average_precision_score, f1_score, precision_score, recall_score
from sklearn.grid_search import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV

from sklearn.svm import SVC
from sklearn.manifold import TSNE

from scipy.stats import expon
import random
```

```

from sklearn.tree import DecisionTreeClassifier

In [16]: pickle_in=open("cleanedData.pickle","rb")
         final = pickle.load(pickle_in)

         '''pickle_in = open("BOW_tfidf_avgW2V_TfidfW2V.pickle","rb")
         count_vect = pickle.load(pickle_in) #BOW
         final_counts = pickle.load(pickle_in) #BOW

         tf_idf_vect = pickle.load(pickle_in) #TFIDF
         final_tf_idf = pickle.load(pickle_in) #TFIDF
         features = pickle.load(pickle_in) #TFIDF

         w2v_model = pickle.load(pickle_in) #w2v
         words = pickle.load(pickle_in) #w2v

         sent_vectors = pickle.load(pickle_in) #avg W2V'''

import pickle
pickle_in = open("BOW_tfidf_avgW2V_Train_test_data.pickle","rb")

count_vect = pickle.load(pickle_in) #BOW
final_counts_train = pickle.load(pickle_in) #BOW
final_counts_test = pickle.load(pickle_in) #BOW
tf_idf_vect = pickle.load(pickle_in) #tfidf
final_tf_idf_train = pickle.load(pickle_in) #tfidf
final_tf_idf_test = pickle.load(pickle_in) #tfidf
features = pickle.load(pickle_in)
sent_vectors_train = pickle.load(pickle_in) #avgW2v Vectors
sent_vectors_test = pickle.load(pickle_in) #avgW2v Vectors

```

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\base.py:315: UserWarning: Trying to unpickle UserWarning)

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\base.py:315: UserWarning: Trying to unpickle UserWarning)

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\base.py:315: UserWarning: Trying to unpickle UserWarning)

```

In [17]: train_data = final.head(int(0.80*final.shape[0]))
         test_data = final.head(int(0.20*final.shape[0])+1)

         scores = final['Score'].get_values()
         len(scores)

```

Out [17]: 364171

Util Functions

```

In [19]: def convScores(scores):
    li = lambda x: 1 if x=='positive' else 0
    final_scores = []
    for i in range(0,len(scores)):
        final_scores.append(li(scores[i]))
    return final_scores

def convToNpArray(arr):
    if(type(arr) == list):
        arr = np.array(arr)
        return arr
    else:
        return arr;

def confusionMatrix(y_test,pred):

    df_cm = pd.DataFrame(confusion_matrix(y_test, pred), index = ['False','True'],
                        columns = ['False','True'])
    sn.heatmap(df_cm, annot=True)
    plt.title('Confusion Matrix')
    plt.ylabel('Actual')
    plt.xlabel('Predicted')
    plt.show()

def show_most_informative_features(vectorizer, clf, n=10):
    feature_names = vectorizer.get_feature_names()
    coefs_with_fns = sorted(zip(clf.feature_importances_, feature_names))
    top = coefs_with_fns[:-(n + 1):-1]
    print("Top 10 Features:")
    for (coef_1, fn_1) in top:
        print("")
        print("%-15s" % (fn_1))

```

## 1.1 Avg w2v DT

```

In [12]: x_1 = sent_vectors_train[0:100000]

# this is only Score/rating of data

y_1 = convScores(train_data['Score'].get_values())[0:100000]

x_test = sent_vectors_test[0:20000]
y_test = convScores(test_data['Score'].get_values())[0:20000]

x_1 = convToNpArray(x_1)
x_test = convToNpArray(x_test)
y_1 = convToNpArray(y_1)

```

```

y_test = convToNpArray(y_test)

In [13]: tuned_parameters = {'max_depth':[11,21,41,61]}

DT_model = DecisionTreeClassifier(class_weight='balanced')
model = GridSearchCV(DT_model,tuned_parameters,
                     scoring='f1',cv=5,n_jobs=-1)

model.fit(x_1,y_1)

print(model.best_estimator_)
print("Score: ",model.score(x_test,y_test))

```

```

DecisionTreeClassifier(class_weight='balanced', criterion='gini',
                       max_depth=61, max_features=None, max_leaf_nodes=None,
                       min_impurity_split=1e-07, min_samples_leaf=1,
                       min_samples_split=2, min_weight_fraction_leaf=0.0,
                       presort=False, random_state=None, splitter='best')
Score: 0.911167653078


```

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid\_search.py:438: ChangedBehaviorWarning: ChangedBehaviorWarning)

```

In [14]: best_DT_model = model.best_estimator_
best_DT_model.fit(x_1,y_1)
pred = best_DT_model.predict(x_test)
confusionMatrix(y_test,pred)


```



```
In [7]: def plotCVvsErr(li_of_depths,x_test,y_test):
        #li_of_depths = [5,7,9,11]
        li_of_errors=[]
        for i in li_of_depths:
            model = DecisionTreeClassifier(max_depth=i,class_weight='balanced')
            model.fit(x_1,y_1)
            err_val = 1-model.score(x_test,y_test);
            li_of_errors.append(err_val)

        plt.plot(li_of_depths,li_of_errors)
        plt.xlabel('Depths')
        plt.ylabel('Error Percentage')
        plt.show()
```

```
In [20]: plotCVvsErr([21,41,121],x_test,y_test)
```



Observation: From the graph we can see that min Error is 0.1 at Depth 41, so best possible value of hyper parameter Depth is 41.

## 1.2 BOW DT

```
In [13]: x_1 = final_counts_train[0:10000]
```

```
# this is only Score/rating of data
```

```
y_1 = convScores(train_data['Score'].get_values())[0:10000]
```

```
x_test = final_counts_test[0:3000]
```

```
y_test = convScores(test_data['Score'].get_values())[0:3000]
```

```
#x_1, x_test, y_1, y_test = train_test_split(x,y, test_size=0.3, random_state=0)
```

```
x_1 = convToNpArray(x_1)
```

```
x_test = convToNpArray(x_test)
```

```
y_1 = convToNpArray(y_1)
```

```
y_test = convToNpArray(y_test)
```

```
In [14]: tuned_parameters = {'max_depth':[151,450,600,900]}
```

```
DT_model = DecisionTreeClassifier(class_weight='balanced')
```

```
model = GridSearchCV(DT_model,tuned_parameters,
                      scoring='f1',cv=5,n_jobs=-1)
```

```
model.fit(x_1,y_1)
```

```
print(model.best_estimator_)
```


```
print("Score: ",model.score(x_test,y_test))
```

```
DecisionTreeClassifier(class_weight='balanced', criterion='gini',
                        max_depth=900, max_features=None, max_leaf_nodes=None,
                        min_impurity_split=1e-07, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        presort=False, random_state=None, splitter='best')
```


Score: 1.0

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid\_search.py:438: ChangedBehaviorWarning: ?  
ChangedBehaviorWarning)

```
In [15]: best_DT_model = model.best_estimator_
best_DT_model.fit(x_1,y_1)
pred = best_DT_model.predict(x_test)
confusionMatrix(y_test,pred)
```



```
In [34]: plotCVvsErr([151,450,600,900],x_test,y_test)
```



### 1.2.1 Important Features

```
In [33]: show_most_informative_features(count_vect,best_DT_model,10)
```

Top 10 Features:

great

best

love

delici

perfect

use

excel

nice

thought



good

### 1.3 TFIDF DT

```
In [8]: x_1 = final_tf_idf_train[0:10000]

        y_1 = convScores(train_data['Score'].get_values())[0:10000]

        x_test = final_tf_idf_test[0:3000]
        y_test = convScores(test_data['Score'].get_values())[0:3000]

        x_1 = convToNpArray(x_1)
        x_test = convToNpArray(x_test)
        y_1 = convToNpArray(y_1)
        y_test = convToNpArray(y_test)
```

```
In [9]: tuned_parameters = {'max_depth':[151,451,651]}

        DT_model = DecisionTreeClassifier(class_weight='balanced')
        model = GridSearchCV(DT_model,tuned_parameters,
                             scoring='f1',cv=5,n_jobs=-1)

        model.fit(x_1,y_1)


        print(model.best_estimator_)
        print("Score: ",model.score(x_test,y_test))
```

```
DecisionTreeClassifier(class_weight='balanced', criterion='gini',
                        max_depth=451, max_features=None, max_leaf_nodes=None,
                        min_impurity_split=1e-07, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        presort=False, random_state=None, splitter='best')
```


Score: 1.0

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid\_search.py:438: ChangedBehaviorWarning: ChangedBehaviorWarning)

```
In [10]: best_DT_model = model.best_estimator_
         best_DT_model.fit(x_1,y_1)
         pred = best_DT_model.predict(x_test)
         confusionMatrix(y_test,pred)
```



In [28]: `plotCVvsErr([121,450,600],x_test,y_test)`



```
In [11]: show_most_informative_features(tf_idf_vect,best_DT_model,10)
```

Top 10 Features:

great

best

love

delici

perfect

good

tea

excel

nice

favorit

## 1.4 Important Features

Observation:

Vectorizer	Depth(Hyper Parameter)
Avg W2v	41
BOW	450
TfIdf	450