

AmazonReviews_KNN_Assignment

May 9, 2018

1 AmazonReviews KNN Assignment

```
In [1]: %matplotlib inline
```

```
import sqlite3
import pandas as pd #for data frames
import numpy as np #numpy array operations
import nltk #natural lang processing, for processing text
import string
import matplotlib.pyplot as plt
import seaborn as sns #for plotting
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
import pickle
import seaborn as sn

import matplotlib.pyplot as plt
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
```

```
C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: 
    "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [2]: pickle_in=open("cleanedData.pickle","rb")
        final = pickle.load(pickle_in)
```

```
In [3]: pickle_in = open("BOW_tfidf_avgW2V_TfidfW2V.pickle","rb")
        count_vect = pickle.load(pickle_in) #BOW
```

```

final_counts = pickle.load(pickle_in) #BOW

tf_idf_vect = pickle.load(pickle_in) #TFIDF
final_tf_idf = pickle.load(pickle_in) #TFIDF
features = pickle.load(pickle_in) #TFIDF

w2v_model = pickle.load(pickle_in) #w2v
words = pickle.load(pickle_in) #w2v

sent_vectors = pickle.load(pickle_in) #avg W2V

```

C:\Users\Dell\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

```

In [4]: # Weighted TF_IDF W2V
        pickle_in = open("WiightedTfidfW2V.pickle", "rb")
        tfidf_sent_vectors = pickle.load(pickle_in)

```

```

In [5]: final.shape

```

```

Out[5]: (364171, 11)

```

```

In [6]: final_counts.shape

```

```

Out[6]: (364171, 115281)

```

```

In [7]: scores = final['Score'].get_values()
        len(scores)

```

```

Out[7]: 364171

```

```

In [8]: li = lambda x: 1 if x=='positive' else 0
        final_scores = []
        for i in range(0,364171):
            final_scores.append(li(scores[i]))

```

1.1 k-fold Cross Validation Using Time Based Split

```

In [57]: from sklearn.model_selection import TimeSeriesSplit

```

```

def k_fold_cross_validation(x_1,y_1):
    li = list(range(0,16))

    #creating odd list of vals for 'k'
    neighbors = list(filter(lambda x: x%2 != 0,li))

    #empty list to hold CV scores/accuracy
    cv_scores = []

```

```

#performing k-fold cross validation with Times Based Split
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k,algorithm='auto',n_jobs=-1)
    tscv = TimeSeriesSplit(n_splits=5)
    acc_sum=0
    cnt=0
    for train_index, test_index in tscv.split(x_1):
        X_train, X_test = x_1[train_index], x_1[test_index]
        Y_train, Y_test = y_1[train_index], y_1[test_index]
        knn.fit(X_train,Y_train)
        pred = knn.predict(X_test)
        acc = accuracy_score(Y_test, pred, normalize=True)*float(100)
        acc_sum=acc_sum+acc
        cnt=cnt+1
    cv_scores.append(acc_sum/cnt)

# determining best k
optimal_k = neighbors[cv_scores.index(max(cv_scores))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, cv_scores)

for xy in zip(neighbors, np.round(cv_scores,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Accuracy Percentage')
plt.show()

print("the Accuracy for each k value is : ", np.round(cv_scores,3))
print('*****')
print("optimal k value is: ",optimal_k)
return optimal_k;

```

2 BOW K-Fold CV

In [58]: final_counts.shape

Out [58]: (364171, 115281)

In [76]: *# Total data frame*

```
x = final_counts[0:10000]
```

```
# this is only Score/rating of data
```

```
y = final_scores[0:10000]
```

```
In [77]: #splitting the train and test data
```

```
x_1, x_test, y_1, y_test = cross_validation.train_test_split(x,y, test_size=0.3, rand
```

```
In [78]: #x_1 = np.array(x_1)
```

```
y_1 = np.array(y_1)
```

```
#x_test = np.array(x_test)
```


```
y_test = np.array(y_test)
```

```
x_1
```

```
Out[78]: <7000x115281 sparse matrix of type '<class 'numpy.int64'>'
        with 365189 stored elements in Compressed Sparse Row format>
```

```
In [79]: optim_k_val = k_fold_cross_validation(x_1,y_1)
```

The optimal number of neighbors is 7.



the Accuracy for each k value is : [80.343 87.993 89.022 89.314 89.177 89.297 89.297 89.297]

optimal k value is: 7

```
In [80]: knn_optimal = KNeighborsClassifier(n_neighbors=optim_k_val)

        # fitting the model
        knn_optimal.fit(x_1, y_1)

        # predict the response
        pred = knn_optimal.predict(x_test)

        # evaluate accuracy
        acc = accuracy_score(y_test, pred) * 100
        print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optim_k_val, acc))
```

The accuracy of the knn classifier for k = 7 is 88.233333%

3 TF_IDF K-fold CV

```
In [81]: final_tf_idf.shape
```

```
Out[81]: (364171, 2910192)
```

```
In [85]: # Total data frame
        x = final_tf_idf[0:10000]

        # this is only Score/rating of data
        y = final_scores[0:10000]
```


```
In [86]: #splitting the train and test data
```

```
x_1, x_test, y_1, y_test = cross_validation.train_test_split(x,y, test_size=0.3, random_state=0)

y_1 = np.array(y_1)
y_test = np.array(y_test)
```

```
In [87]: opt_val = k_fold_cross_validation(x_1,y_1)
```

The optimal number of neighbors is 7.



the Accuracy for each k value is : [85.832 89.503 89.605 89.777 89.623 89.64 89.605 89.503]

 optimal k value is: 7

```
In [88]: knn_optimal = KNeighborsClassifier(n_neighbors=opt_val)

# fitting the model
knn_optimal.fit(x_1, y_1)

# predict the response
pred = knn_optimal.predict(x_test)

# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (opt_val, acc))
```

The accuracy of the knn classifier for k = 7 is 88.766667%

4 W2V K-Fold CV

```
In [17]: # Total data frame
x = sent_vectors[0:50000]
```

```
# this is only Score/rating of data
y = final_scores[0:50000]
```

```
In [18]: #splitting the train and test data
```


```
x_1, x_test, y_1, y_test = cross_validation.train_test_split(x,y, test_size=0.2, random_state=42)

y_1 = np.array(y_1)
y_test = np.array(y_test)

x_1 = np.array(x_1)
x_test = np.array(x_test)
```

```
In [19]: opt_val = k_fold_cross_validation(x_1,y_1)
```

The optimal number of neighbors is 15.



the Accuracy for each k value is : [80.297 86.106 87.888 88.431 88.614 88.665 88.704 88.719 88.719 88.719 88.719 88.719 88.719 88.719 88.719]

```
*****
optimal k value is: 15
```

```

In [20]: knn_optimal = KNeighborsClassifier(n_neighbors=opt_val)

         # fitting the model
         knn_optimal.fit(x_1, y_1)

         prediction = knn_optimal.predict(x_test)

         acc = accuracy_score(y_test, prediction) * 100

         print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (opt_val, acc))

```

The accuracy of the knn classifier for k = 15 is 88.700000%

5 Weighted TF-IDF W2V K-Fold CV

```

In [89]: len(tfidf_sent_vectors)

```

```

Out[89]: 1000

```

```

In [94]: x = tfidf_sent_vectors[0:300]

```

```

         y = final_scores[0:300]

```

```

In [95]: x_1, x_test, y_1, y_test = cross_validation.train_test_split(x,y, test_size=0.3, random_state=42)

```

```

         y_1 = np.array(y_1)
         y_test = np.array(y_test)

```

```

         x_1 = np.array(x_1)
         x_test = np.array(x_test)

```

```

In [96]: len(y_test)

```

```

Out[96]: 90

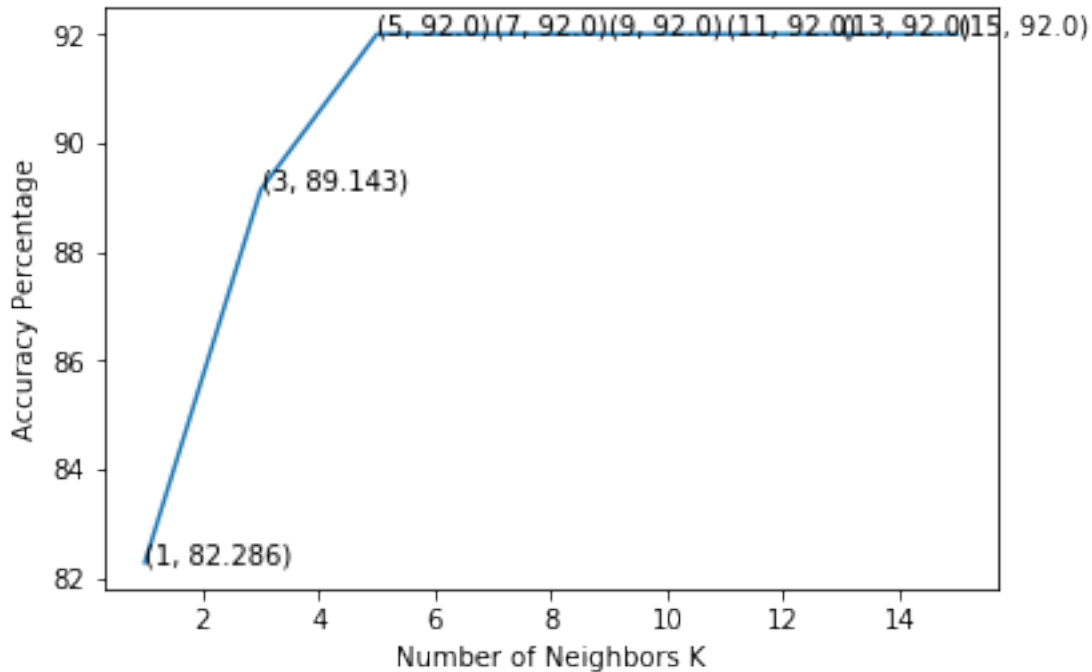
```

```

In [97]: opt_val = k_fold_cross_validation(x_1,y_1)

```

The optimal number of neighbors is 5.



the Accuracy for each k value is : [82.286 89.143 92. 92. 92. 92. 92. 92.]

 optimal k value is: 5

```
In [98]: knn_optimal = KNeighborsClassifier(n_neighbors=3)

knn_optimal.fit(x_1,y_1)

pred = knn_optimal.predict(x_test)

acc = accuracy_score(y_test,pred) * 100

print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (opt_val, acc))
```

The accuracy of the knn classifier for k = 5 is 97.777778%

6 Summary:

After performing K-Fold Cross Validation here are the results of diff techniques:

- 1) BOW -> k = 7 acc= 88.23%

- 2) TF-IDF -> $k = 7$ acc = 88.7%
- 3) W2V -> $k = 15$ acc = 88.7%
- 4) Weighted TF-IDF W2V -> $k = 5$ acc = 97.7%(due to memory overflow used less datapoints)