

AmazonReviews_RF_GBDT

September 2, 2018

1 Amazon Reviews Classification Using Random Forest and Gradient-Boosting

```
In [10]: %matplotlib inline
```

```
import pandas as pd #for data frames
import numpy as np #numpy array operations
import string
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
import pickle
from sklearn.cross_validation import train_test_split
import seaborn as sn
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from sklearn import cross_validation
from sklearn.model_selection import train_test_split
from sklearn.metrics import average_precision_score, f1_score, precision_score, recall_score
from sklearn.grid_search import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import roc_curve, auc

from scipy.stats import expon
import random

from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
In [2]: pickle_in=open("cleanedData.pickle","rb")
        final = pickle.load(pickle_in)
```

```
import pickle
pickle_in = open("BOW_tfidf_avgW2V_Train_test_data.pickle","rb")
```

```

count_vect = pickle.load(pickle_in) #BOW
final_counts_train = pickle.load(pickle_in) #BOW
final_counts_test = pickle.load(pickle_in) #BOW
tf_idf_vect = pickle.load(pickle_in) #tfidf
final_tf_idf_train = pickle.load(pickle_in) #tfidf
final_tf_idf_test = pickle.load(pickle_in) #tfidf
features = pickle.load(pickle_in)
sent_vectors_train = pickle.load(pickle_in) #avgW2v Vectors
sent_vectors_test = pickle.load(pickle_in) #avgW2v Vectors

```

```

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\base.py:315: UserWarning: Trying to unpickle
UserWarning)
C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\base.py:315: UserWarning: Trying to unpickle
UserWarning)
C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\base.py:315: UserWarning: Trying to unpickle
UserWarning)

```

```

In [3]: train_data = final.head(int(0.80*final.shape[0]))
        test_data = final.head(int(0.20*final.shape[0])+1)

```

```

scores = final['Score'].get_values()
len(scores)

```

```

Out[3]: 364171

```

```

In [28]: def convScores(scores):
        li = lambda x: 1 if x=='positive' else 0
        final_scores = []
        for i in range(0,len(scores)):
            final_scores.append(li(scores[i]))
        return final_scores

    def convToNpArray(arr):
        if(type(arr) == list):
            arr = np.array(arr)
            return arr
        else:
            return arr;

    def confusionMatrix(y_test,pred):
        df_cm = pd.DataFrame(confusion_matrix(y_test, pred), index = ['False','True'],
                               columns = ['False','True'])
        sn.heatmap(df_cm, annot=True)
        plt.title('Confusion Matrix')
        plt.ylabel('Actual')
        plt.xlabel('Predicted')
        plt.show()

    def auc_roc(y_test,pred):
        fpr, tpr, thresholds = roc_curve(y_test,pred)

```

```

acc = auc(fpr, tpr)
print("Area Under The Curve is : ",acc)
plt.figure()

plt.plot(fpr, tpr, color='darkorange',
         label='ROC curve (area/auc = %0.2f)' % acc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic Curve')
plt.legend(loc="lower right")
plt.show()
def show_most_informative_features(vectorizer, clf, n=20):
    feature_names = vectorizer.get_feature_names()
    coefs_with_fns = sorted(zip(clf.feature_importances_, feature_names))
    top = zip(coefs_with_fns[:n], coefs_with_fns[:-(n + 1):-1])
    print("These are the top {} important Features Which are most widely used in Posi")
    print("")
    print("\tPositive: \t\t\tNegative:")
    for (coef_1, fn_1), (coef_2, fn_2) in top:
        print("")
        print("\t%-15s\t\t\t\t%-15s" % (fn_2,fn_1))

```

2 RF

2.1 Avg W2V RF

In [29]: `x_1 = sent_vectors_train[0:50000]`

```

# this is only Score/rating of data

y_1 = convScores(train_data['Score'].get_values())[0:50000]

x_test = sent_vectors_test[0:6000]
y_test = convScores(test_data['Score'].get_values())[0:6000]

x_1 = convToNpArray(x_1)
x_test = convToNpArray(x_test)
y_1 = convToNpArray(y_1)
y_test = convToNpArray(y_test)

```

Considering `n_estimators/k` and `max_depth` as hyperparameters

In [30]: `tuned_parameters = {'n_estimators': [10,20,40,50],`
`'max_depth': [20,50,100]}`

```

rf_model = RandomForestClassifier(n_jobs=-1,oob_score=True,class_weight='balanced')
model = GridSearchCV(rf_model,tuned_parameters,
                    scoring='accuracy',cv=3,n_jobs=-1)

model.fit(x_1,y_1)

print(model.best_estimator_)
print("Score: ",model.score(x_test,y_test))

```

```

RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=20, max_features='auto',
                        max_leaf_nodes=None, min_impurity_split=1e-07,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=40, n_jobs=-1,
                        oob_score=True, random_state=None, verbose=0, warm_start=False)
Score: 0.882833333333


```

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning: ChangedBehaviorWarning)

```

In [31]: best_rf_model = model.best_estimator_
best_rf_model.fit(x_1,y_1)
pred = best_rf_model.predict(x_test)
confusionMatrix(y_test,pred)

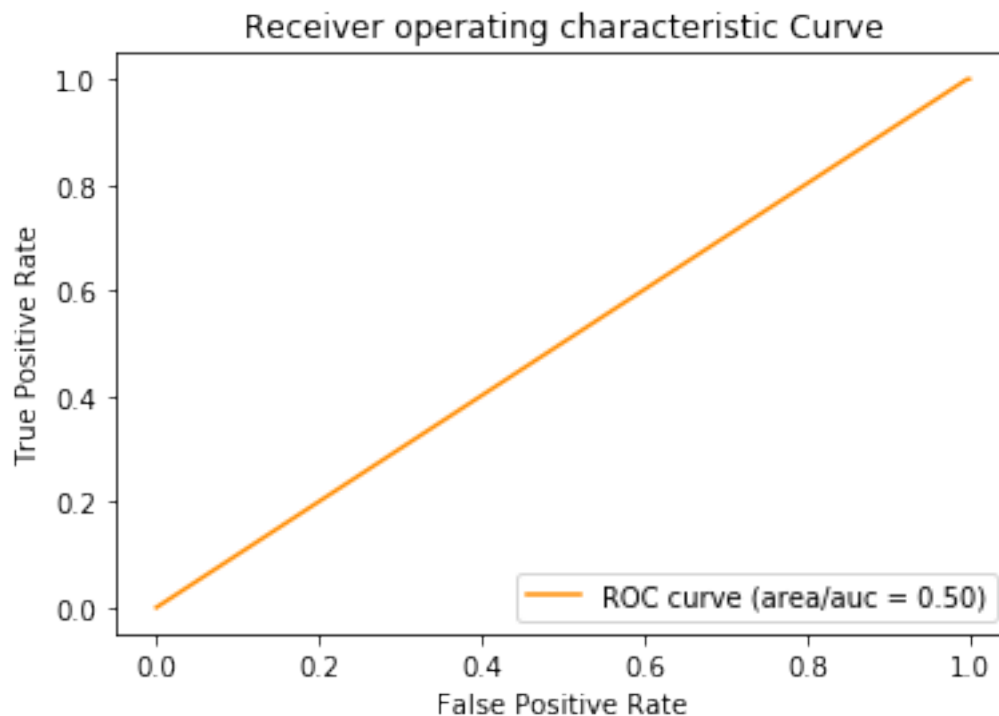
```



AUC

```
In [16]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.50131996277



2.2 BOW RF

```
In [33]: # Total data frame
```

```
x_1 = final_counts_train[0:50000]
```

```
# this is only Score/rating of data
```

```
y_1 = convScores(train_data['Score'].get_values())[0:50000]
```

```
x_test = final_counts_test[0:6000]
```

```
y_test = convScores(test_data['Score'].get_values())[0:6000]
```

```
#x_1, x_test, y_1, y_test = train_test_split(x,y, test_size=0.3, random_state=0)
```

```

x_1 = convToNpArray(x_1)
x_test = convToNpArray(x_test)
y_1 = convToNpArray(y_1)
y_test = convToNpArray(y_test)

```

```

In [34]: tuned_parameters = {'n_estimators':[10,20,40,50],
                             'max_depth':[20,50,100]}

```

```

rf_model = RandomForestClassifier(n_jobs=-1,oob_score=True,class_weight='balanced')
model = GridSearchCV(rf_model,tuned_parameters,
                     scoring='accuracy',cv=3,n_jobs=-1)

model.fit(x_1,y_1)

print(model.best_estimator_)
print("Score: ",model.score(x_test,y_test))

```

```

RandomForestClassifier(bootstrap=True, class_weight='balanced',
                       criterion='gini', max_depth=50, max_features='auto',
                       max_leaf_nodes=None, min_impurity_split=1e-07,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, n_estimators=50, n_jobs=-1,
                       oob_score=True, random_state=None, verbose=0, warm_start=False)

```


Score: 0.974

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning: ?
 ChangedBehaviorWarning)

```

In [35]: best_rf_model = model.best_estimator_
         best_rf_model.fit(x_1,y_1)
         pred = best_rf_model.predict(x_test)
         confusionMatrix(y_test,pred)


```



AUC

```
In [21]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.940581366171



```
In [37]: show_most_informative_features(count_vect,best_rf_model,10)
```

These are the top 10 important Features Which are most widely used in Positive and Negative Reviews

Positive:	Negative:
great	aa
love	aaa
disappoint	aaaa
best	aaaaa
delici	aaaaaaaaaaaaa
would	aaa
perfect	aaaaaaaaaaaaaaaaaaaaargh
wast	aaaaaaaaagghh
favorit	aaaaaaahhhhhh

worst

aaaaaaarrrrrrggghhh

2.3 TFidf

```
In [22]: x_1 = final_tf_idf_train[0:10000]
```

```
y_1 = convScores(train_data['Score'].get_values())[0:10000]
```

```
x_test = final_tf_idf_test[0:3000]
```

```
y_test = convScores(test_data['Score'].get_values())[0:3000]
```

```
x_1 = convToNpArray(x_1)
```

```
x_test = convToNpArray(x_test)
```

```
y_1 = convToNpArray(y_1)
```

```
y_test = convToNpArray(y_test)
```

```
In [23]: tuned_parameters = {'n_estimators': [10, 20, 40, 50],  
                             'max_depth': [20, 50, 100]}
```

```
rf_model = RandomForestClassifier(n_jobs=-1, oob_score=True, class_weight='balanced')
```

```
model = GridSearchCV(rf_model, tuned_parameters,
```

```
                      scoring='accuracy', cv=3, n_jobs=-1)
```

```
model.fit(x_1, y_1)
```

```
print(model.best_estimator_)
```


```
print("Score: ", model.score(x_test, y_test))
```

```
RandomForestClassifier(bootstrap=True, class_weight='balanced',  
                        criterion='gini', max_depth=50, max_features='auto',  
                        max_leaf_nodes=None, min_impurity_split=1e-07,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, n_estimators=50, n_jobs=-1,  
                        oob_score=True, random_state=None, verbose=0, warm_start=False)
```

```
C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning:   
  ChangedBehaviorWarning)
```


```
Score: 0.982
```

```
In [24]: best_rf_model = model.best_estimator_  
best_rf_model.fit(x_1, y_1)  
pred = best_rf_model.predict(x_test)  
confusionMatrix(y_test, pred)
```



```
In [25]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.914667289577



2.4 Important Features

In [27]: `show_most_informative_features(tf_idf_vect,best_rf_model,10)`

These are the top 10important Features Which are most widely used in Positive and Negative Rev.

Positive:

money

order

return

great

horribl

away

snack

love

Negative:

aa

aa pleas

aa state

aaa

aaa aaa

aaa condit

aaa dont

aaa hockey

high recommend	aaa job
bland	aaa magazin

3 GBDT

3.1 Avg W2V

```
In [48]: x_1 = sent_vectors_train[0:30000]

# this is only Score/rating of data

y_1 = convScores(train_data['Score'].get_values())[0:30000]

x_test = sent_vectors_test[0:6000]
y_test = convScores(test_data['Score'].get_values())[0:6000]

x_1 = convToNpArray(x_1)
x_test = convToNpArray(x_test)
y_1 = convToNpArray(y_1)
y_test = convToNpArray(y_test)

In [49]: tuned_parameters = {'n_estimators': [50, 75, 100, 125],
                             'learning_rate': [0.01, 0.1, 0.5],
                             'max_depth': [3, 5, 7]}

#using 70% of data to train the base models (Row Sampling)
gb_model = GradientBoostingClassifier(subsample=0.7)
model = GridSearchCV(gb_model, tuned_parameters,
                    scoring='accuracy', cv=3, n_jobs=-1)

model.fit(x_1, y_1)


print(model.best_estimator_)
print("Score: ", model.score(x_test, y_test))

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                          learning_rate=0.1, loss='deviance', max_depth=7,
                          max_features=None, max_leaf_nodes=None,
                          min_impurity_split=1e-07, min_samples_leaf=1,
                          min_samples_split=2, min_weight_fraction_leaf=0.0,
                          n_estimators=100, presort='auto', random_state=None,
                          subsample=0.7, verbose=0, warm_start=False)

Score: 0.887666666667
```


C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning: ChangedBehaviorWarning)

```
In [50]: best_gb_model = model.best_estimator_  
best_gb_model.fit(x_1,y_1)  
pred = best_gb_model.predict(x_test)  
confusionMatrix(y_test,pred)
```



```
In [51]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.539169985391



3.1.1 Using XGBClassifier

```
In [54]: tuned_parameters = {'n_estimators': [75, 100, 120],
                             'learning_rate': [0.01, 0.1, 0.5]}

#using 70% of data to train the base models (Row Sampling)
gb_model = XGBClassifier(subsample=0.7, n_jobs=4)
model = GridSearchCV(gb_model, tuned_parameters,
                     scoring='accuracy', cv=3, n_jobs=-1)


model.fit(x_1, y_1)

print(model.best_estimator_)
print("Score: ", model.score(x_test, y_test))

XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
               max_depth=3, min_child_weight=1, missing=None, n_estimators=120,
               n_jobs=4, nthread=None, objective='binary:logistic', random_state=0,
               reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
               silent=True, subsample=0.7)
Score: 0.883333333333
```


C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning:
 ChangedBehaviorWarning)

```
In [55]: best_gb_model = model.best_estimator_  
best_gb_model.fit(x_1,y_1)  
pred = best_gb_model.predict(x_test)  
confusionMatrix(y_test,pred)
```



```
In [56]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.517818228073



4 BOW

In [38]: *# Total data frame*

```
x_1 = final_counts_train[0:30000]
```

```
# this is only Score/rating of data
```

```
y_1 = convScores(train_data['Score'].get_values())[0:30000]
```

```
x_test = final_counts_test[0:6000]
```

```
y_test = convScores(test_data['Score'].get_values())[0:6000]
```

```
#x_1, x_test, y_1, y_test = train_test_split(x,y, test_size=0.3, random_state=0)
```

```
x_1 = convToNpArray(x_1)
```

```
x_test = convToNpArray(x_test)
```

```
y_1 = convToNpArray(y_1)
```

```
y_test = convToNpArray(y_test)
```

In [39]: `tuned_parameters = {'n_estimators': [75, 100, 120],
 'learning_rate': [0.01, 0.1, 0.5]}`


```

#using 70% of data to train the base models (Row Sampling)
gb_model = XGBClassifier(subsample=0.7,n_jobs=4)
model = GridSearchCV(gb_model,tuned_parameters,
                     scoring='accuracy',cv=3,n_jobs=-1)

model.fit(x_1,y_1)

print(model.best_estimator_)
print("Score: ",model.score(x_test,y_test))

XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bytree=1, gamma=0, learning_rate=0.5, max_delta_step=0,
              max_depth=3, min_child_weight=1, missing=None, n_estimators=120,
              n_jobs=4, nthread=None, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=True, subsample=0.7)
Score:  0.9375

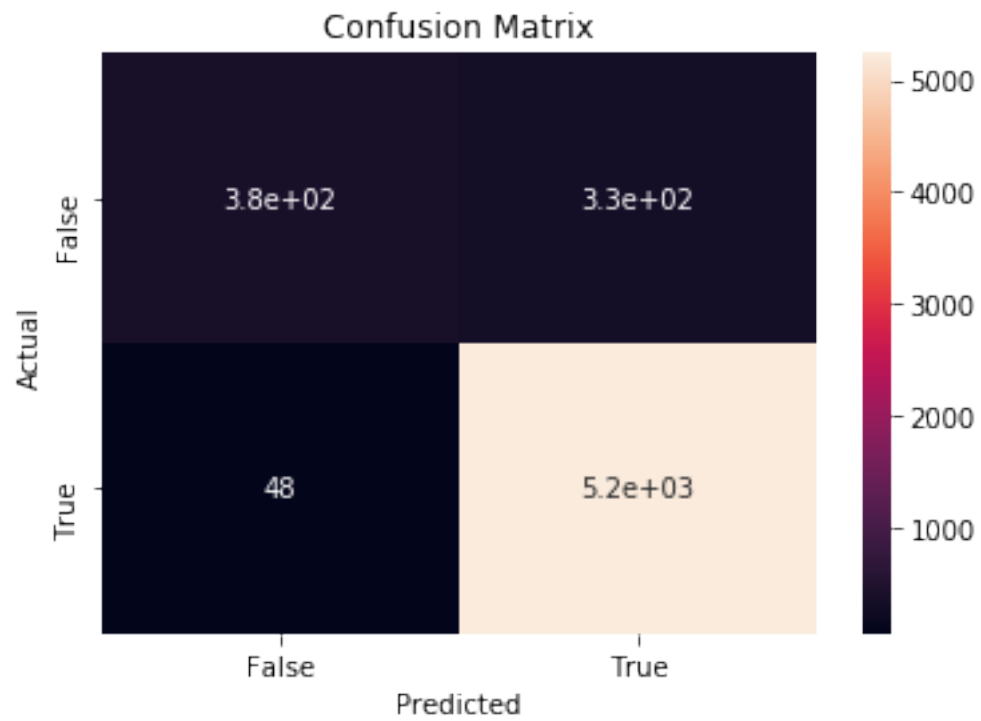
```

C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning: ChangedBehaviorWarning)

```


In [40]: best_gb_model = model.best_estimator_
         best_gb_model.fit(x_1,y_1)
         pred = best_gb_model.predict(x_test)
         confusionMatrix(y_test,pred)

```



```
In [41]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.764206869258



4.1 Important Features

In [42]: `show_most_informative_features(count_vect,best_gb_model,10)`

These are the top 10 important Features Which are most widely used in Positive and Negative Reviews

Positive:

Negative:

like

aa

great

aaa

tast

aaaa

product

aaaaa

love

aaaaaaaaaaaaaa

best

aa

tri

aaaaaaaaaaaaaaaaaaaaargh

littl

aaaaaaaaagghh

good	aaaaaaahhhhhh
flavor	aaaaaaarrrrrggghhh

4.2 TFidf

```
In [43]: x_1 = final_tf_idf_train[0:10000]
```

```
y_1 = convScores(train_data['Score'].get_values())[0:10000]
```

```
x_test = final_tf_idf_test[0:3000]
```

```
y_test = convScores(test_data['Score'].get_values())[0:3000]
```

```
x_1 = convToNpArray(x_1)
```

```
x_test = convToNpArray(x_test)
```

```
y_1 = convToNpArray(y_1)
```

```
y_test = convToNpArray(y_test)
```

```
In [44]: tuned_parameters = {'n_estimators': [75, 100, 120],
                             'learning_rate': [0.01, 0.1, 0.5]}
```

```
#using 70% of data to train the base models (Row Sampling)
```

```
gb_model = XGBClassifier(subsample=0.7, n_jobs=4)
```

```
model = GridSearchCV(gb_model, tuned_parameters,
                     scoring='accuracy', cv=3, n_jobs=-1)
```

```
model.fit(x_1, y_1)
```

```
print(model.best_estimator_)
```

```
print("Score: ", model.score(x_test, y_test))
```


```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bytree=1, gamma=0, learning_rate=0.5, max_delta_step=0,
              max_depth=3, min_child_weight=1, missing=None, n_estimators=120,
              n_jobs=4, nthread=None, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=True, subsample=0.7)
```

```
C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\grid_search.py:438: ChangedBehaviorWarning:
  ChangedBehaviorWarning)
```

```
Score: 0.961666666667
```


```
In [45]: best_gb_model = model.best_estimator_
         best_gb_model.fit(x_1, y_1)
```

```
pred = best_gb_model.predict(x_test)
confusionMatrix(y_test,pred)
```



```
In [46]: auc_roc(y_test,pred)
```

Area Under The Curve is : 0.847486197897



4.3 Important Features

In [47]: `show_most_informative_features(tf_idf_vect,best_gb_model,10)`

These are the top 10 important Features Which are most widely used in Positive and Negative Reviews

Positive:

good

like

one

tast

tri

product

get

use

Negative:

aa

aa pleas

aa state

aaa

aaa aaa

aaa condit

aaa dont

aaa hockey

love

aaa job

great

aaa magazin

Summary:

Classifier	Vectorizer	Hyper Parameters	Confusion Matrix	Accuracy	AUC
Random Forest	Avg W2v	n_estimators: 40 max_depth: 20	TPR :0.999622 TNR : 0.004243 FPR : 0.995757 FNR: 0.000378	88%	0.5
Random Forest	BOW	n_estimators: 40 max_depth: 50	TPR :0.983374 TNR : 0.903819 FPR : 0.096181 FNR: 0.016626	97.5%	0.94
Random Forest	TfIdf	n_estimators: 50			

max_depth: 50</td>	
<td> TPR :0.999625 TNR : 0.847305 FPR : 0.152695 FNR: 0.000375 </td>	
<td>97.9%</td>	
<td>0.91</td>	

<td>Gradient Boost / XGBClassifier</td>	
<td>Avg W2v</td>	
<td> n_estimators: 100 learning_rate: 0.1 max_depth: 7 Using XGBClassifier n_estimators: 120 learning_rate: 0.1 </td>	
<td> TPR :0.983374 TNR : 0.128713 FPR : 0.871287 FNR: 0.016626 </td>	
<td>88.33%</td>	
<td>0.54</td>	

<td>Gradient Boost / XGBClassifier</td>	
<td>BOW</td>	
<td>n_estimators: 120 learning_rate: 0.5</td>	
<td> TPR :0.990931 TNR : 0.537482 FPR : 0.462518 FNR: 0.009069 </td>	
<td>93.75%</td>	
<td>0.76</td>	

<td>Gradient Boost / XGBClassifier</td>	
<td>TfIdf</td>	
<td>n_estimators: 120 learning_rate: 0.5 </td>	
<td> TPR :0.994374 TNR : 0.700599 FPR : 0.299401 FNR: 0.005626	

