

SGD_on_LinearRegression

September 1, 2018

1 Linear Regression using Stochastic Gradient Descent(SGD) on Boston House Prices

```
In [90]: from sklearn.datasets import load_boston
import numpy as np
import random
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.utils import shuffle
from sklearn.metrics import r2_score
from datetime import datetime
from sklearn.linear_model import SGDRegressor
from matplotlib.pyplot import scatter
```

```
In [3]: boston_data = load_boston()
```

```
In [4]: X = boston_data.data
```

```
In [5]: X.shape
```

```
Out[5]: (506, 13)
```

Observation: There are 506 Data points with 13 Features

```
In [6]: y = boston_data.target
len(y)
```

```
Out[6]: 506
```

```
In [7]: from sklearn.cross_validation import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X,y,test_size=0.33)
```

```
C:\Users\Dell\Anaconda3\lib\site-packages\sklearn\cross_validation.py:44: DeprecationWarning: 
    "This module will be removed in 0.20.", DeprecationWarning)
```

1.0.1 Standardization

```
In [8]: std = StandardScaler()
        x_train = std.fit_transform(x_train)
        x_test = std.fit_transform(x_test)
```

1.1 Custom Linear Regression using Gradient Descent

```
In [9]: def linearRegressionGD(X,y,w = np.zeros(13),b = 0,iterations=5000,learning_rate=0.1,epsilon=0.0001):
        """
            Linear Regression Using Gradient Descent Optimization

            Parameters
            -----
            X : numpy array or sparse matrix of shape [n_samples,n_features]
                Training data

            y : numpy array of shape [n_samples, n_targets]
                Target values

            w: Weight Vector first initialized with 0's

            b: y-intercept

            iterations: no.of iterations

            epsilon_val: Termination Criteria

            Returns
            -----
            w,b,cost or MSE
        """
        start = datetime.now()
        n = float(len(y))
        costs = []
        decay_rate = 0.0001
        eta=learning_rate
        for i in range(iterations):

            #learning rate
            eta = eta * (1.0/(1+decay_rate*i))
            error = y - predict(X,w,b) #y_act - y_actual

            #partial differentiation w.r.t w
            w_grad = (-2.0/n)*(X.T.dot(error))
```

```

        #partial differentiation w.r.t b
        b_grad = (-2.0/n)*(np.sum(error))

        #MSE
        cost = (1.0/n)*np.sum(np.power(error,2))

        #updating weight vector
        w = w - (eta*w_grad)

        #updating y-intercept
        b = b - (eta*b_grad)

        costs.append(cost)

        #Stopping Criteria
        if i==0:
            w_prev = w;
        else:
            dist = (np.linalg.norm(w-w_prev))
            w_prev = w

            #Terminates the loop when difference between w and w_prev is very less
            if(round(dist,6)<epsilon_val):
                print("no.of iterations: ",i,"Cost: ",min(costs))
                break;

    print("Time: Taken: ",datetime.now() - start)
    return w,b,cost;

def predict(x,m,c):
    """
        It predicts the target value based on the input parameter values

        Parameters
        -----
        x : numpy array or sparse matrix of shape [n_samples,n_features]
            Training data

    """
    y = np.dot(x,m) + c
    return y

def score(X,y,m,c):

```

```

"""
It gives the Score or accuracy

Parameters
-----
X : numpy array or sparse matrix of shape [n_samples,n_features]
    Training data

y : numpy array of shape [n_samples, n_targets]
    Target values

m : slope

c : y-intercept
"""

from sklearn.metrics import r2_score
return r2_score(y,predict(X,m,c).flatten())

```

```
In [20]: w1,b1,cost1 = linearRegressionGD(x_train,y_train)
```

```
no.of iterations: 131 Cost: 20.9300479465
Time: Taken: 0:00:00.011941
```

```
In [21]: score(x_test,y_test,w1,b1)
```

```
Out[21]: 0.7368587800798797
```

1.2 Custom Linear Regression using SGD

```
In [56]: def linearRegressionSGD(X,y,w = np.zeros(13),b = 0,iterations=1000,case_size=100,learn
"""
    Linear Regression Using Gradient Descent Optimization

Parameters
-----
X : numpy array or sparse matrix of shape [n_samples,n_features]
    Training data

y : numpy array of shape [n_samples, n_targets]
    Target values

w: Weight Vector first initialized with 0's

b: y-intercept

```

iterations: no.of iterations

case_size : bunch of points to be taken in each iteration

epsilon_val: Termination Criteria

Returns

w,b, cost or MSE

"""

```
start = datetime.now()
n = float(len(y))
costs = []
decay_rate = 0.0001
eta=learning_rate

for iters in range(iterations):
    eta = eta * (1.0/(1+decay_rate*iters))

    X,y = shuffle(X,y,random_state=0,replace=True)

    for i in range(case_size):
        x_mini = X[i:i+case_size]
        y_mini = y[i:i+case_size]

        error = y_mini - predict(x_mini,w,b) #y_act - y_predicted

        #partial differentiation w.r.t w
        w_grad = (-2.0/n)*(x_mini.T.dot(error))

        #partial differentiation w.r.t b
        b_grad = (-2.0/n)*(np.sum(error))

        #MSE
        cost = (1.0/n)*np.sum(np.power(error,2))

        #updating weight vector
        w = w - (eta*w_grad)

        #updating y-intercept
        b = b - (eta*b_grad)
```

```

        costs.append(cost)

    #Stopping Criteria
    if iters==0:
        w_prev = w;
    else:
        dist = (np.linalg.norm(w-w_prev))
        w_prev = w

    #Terminates the loop when difference between w and w_prev is very less
    if(round(dist,6)<epsilon_val):
        print("no.of iterations: ",iters,"Cost: ",min(costs))
        break;

    print("Time Taken: ",datetime.now() - start)
    return w,b,cost;

```

```
In [57]: w,b,cost = linearRegressionSGD(x_train,y_train)
```

```
no.of iterations: 379 Cost: 2.3610421482
```

```
Time Taken: 0:00:00.901566
```

```
In [58]: # final Weight Vector
         w
```

```
Out[58]: array([-0.44528049,  1.18336573, -0.00825891,  0.71232241, -2.15167063,
                2.69403028, -0.35533995, -3.51291847,  2.14806818, -1.61030112,
                -2.13751993,  0.96277588, -3.32768966])
```

```
In [59]: # Final y-intercept
         b
```

```
Out[59]: 22.598247491871987
```

```
In [60]: cost
```

```
Out[60]: 6.0160838332641928
```

```
In [61]: score(x_test,y_test,w,b)
```

```
Out[61]: 0.7428985195078156
```

1.3 Sklearn SGD Regressor

```
In [54]: sgd_lr = SGDRegressor(eta0=0.1,shuffle=True)
```

```
sgd_lr.fit(x_train,y_train)
```

```
y_pred = sgd_lr.predict(x_test)

r2_score(y_test,y_pred)
```

```
Out [54]: 0.7236731566748813
```

```
In [55]: sgd_lr.coef_
```

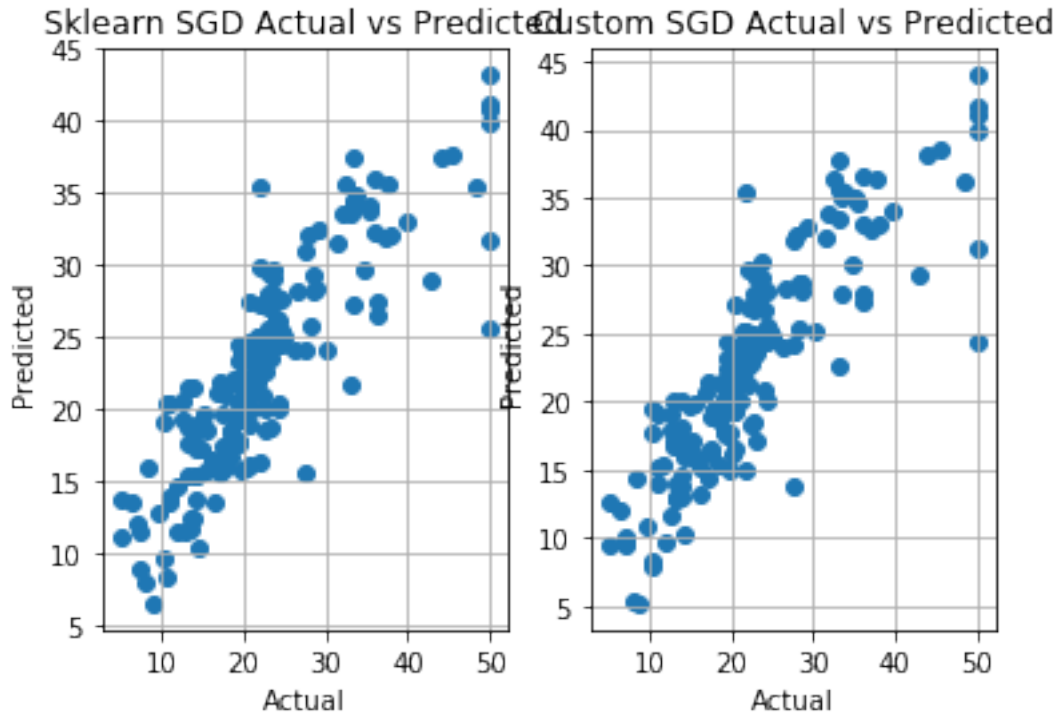
```
Out [55]: array([-0.51154757,  1.16277626,  0.45075066,  0.54998761, -2.0917407 ,
                2.51920819, -0.59341105, -3.59284535,  1.99623011, -1.07285643,
                -2.19643933,  1.53888722, -3.21550559])
```

1.4 Scatter Plots of Actual Vs Predictions of Sklearn and Custom Impls

```
In [91]: plt.subplot(121)
plt.scatter(x=y_test, y=y_pred);
plt.title('Sklearn SGD Actual vs Predicted')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.grid()

plt.subplot(122)
plt.scatter(x=y_test, y=predict(x_test,w,b).flatten());
plt.title('Custom SGD Actual vs Predicted')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.grid()

plt.show()
```



1.5 Weight Vectors Comparison

```
In [64]: print("Custom GD Linear Regression Weight Vector: ",w1)
         print("Custom SGD Linear Regression Weight Vector: ",w)
         print("SkLearn SGD Linear Regression Weight Vector: ",sgd_lr.coef_)
```

```
Custom GD Linear Regression Weight Vector:  [-0.31488545  1.01598113 -0.04944872  0.77237078
 -0.49479217 -3.30437067  1.55003278 -1.20937898 -2.08453311  0.99504223
 -3.0701105 ]
```

```
Custom SGD Linear Regression Weight Vector:  [-0.44528049  1.18336573 -0.00825891  0.71232241
 -0.35533995 -3.51291847  2.14806818 -1.61030112 -2.13751993  0.96277588
 -3.32768966]
```

```
SkLearn SGD Linear Regression Weight Vector:  [-0.51154757  1.16277626  0.45075066  0.5499876
 -0.59341105 -3.59284535  1.99623011 -1.07285643 -2.19643933  1.53888722
 -3.21550559]
```

1.6 Scores Custom vs Sklearn Implementation

```
In [63]: print("Custom GD Linear Regression Score: ",score(x_test,y_test,w1,b1))
         print("Custom SGD Linear Regression Score: ",score(x_test,y_test,w,b))
         print("SkLearn SGD Linear Regression Score: ",r2_score(y_test,y_pred))
```

```
Custom GD Linear Regression Score:  0.73685878008
```

```
Custom SGD Linear Regression Score:  0.742898519508
```


SkLearn SGD Linear Regression Score: 0.723673156675

Accuracy is almost similar, Custome Sgd Implementation got highest score here, but Scores may change slightly based on the shuffle.